

RECIPE

Hardware-Accelerated Replication Protocols

Dimitra Giantsidi, **Emmanouil (Manos) Giortamis**, Julian Pritzi,
Maurice Bailleu, Manos Kapritsos, Pramod Bhatotia



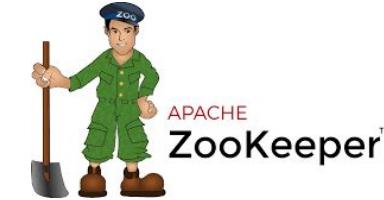
THE UNIVERSITY
of EDINBURGH



Technische
Universität
München

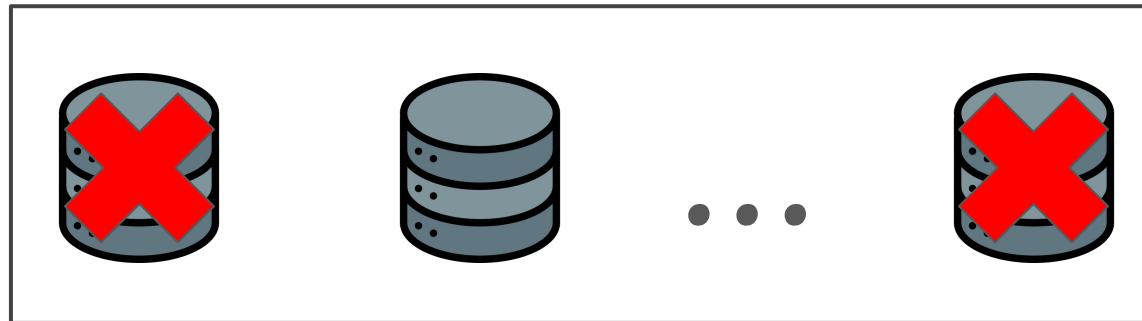


Distributed systems power everything



Distributed systems are the foundation of modern cloud infrastructure

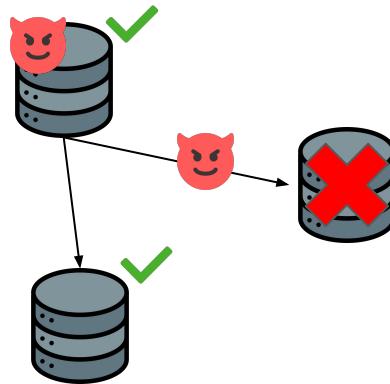
Distributed systems are prone to failures



How to make distributed systems fault-tolerant?

Fault tolerance models

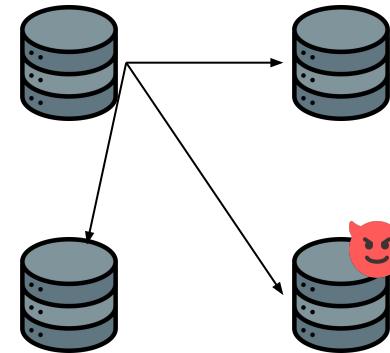
Crash Fault Tolerance (CFT)



CFT system

$2f+1$ nodes handle f failures

Byzantine Fault Tolerance (BFT)



BFT system

$3f+1$ nodes handle f failures

Which fault tolerance model is the best?

Crash Fault Tolerance (CFT)

- High performance
- Simplicity
- Vulnerability



Byzantine Fault Tolerance (BFT)

- Robustness
- Complexity
- Poor performance

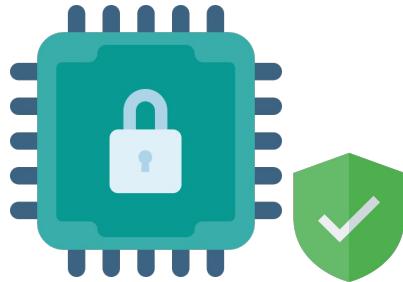


In the modern untrusted cloud we need BFT guarantees, but they are expensive

Research question

How do we systematically design trustworthy distributed systems for Byzantine cloud environments while offering high performance and scalability?

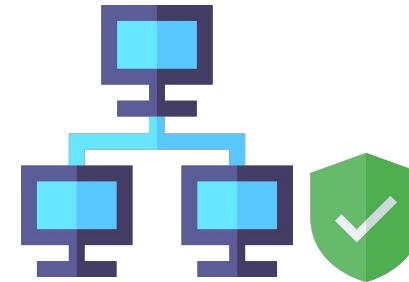
Key insight: Leverage modern hardware for BFT



Trusted computing



BFT robustness



Userspace networking



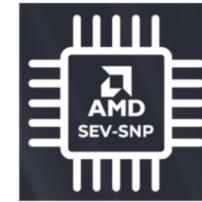
Performance

We can have CFT simplicity and performance with BFT robustness!

Trusted computing for BFT robustness



- CPU-based Trusted Execution Environments (TEEs)
- TEEs provide hardware isolation → protocol compliance
- BFT with TEEs requires $2f+1$ nodes, same as CFT!



Trusted computing can make BFT systems **scalable**

Userspace I/O for BFT systems



- High-throughput and low-latency networking
- Kernel bypass → less system calls
- High performance for TEE-based systems



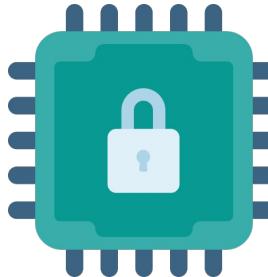
Userspace networking can make BFT systems **performant**

The complete RECIPE



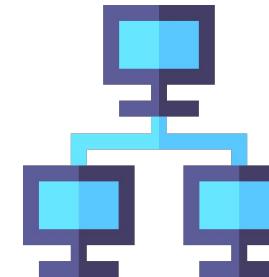
CFT Protocol

+



TEEs

+



Userspace I/O

- Simplicity 

- Robustness 

- Performance 

RECIPE: Hardware-Accelerated Replication Protocols

Rethinking Crash Fault Tolerance Protocols for Untrusted Cloud Environments

Properties:

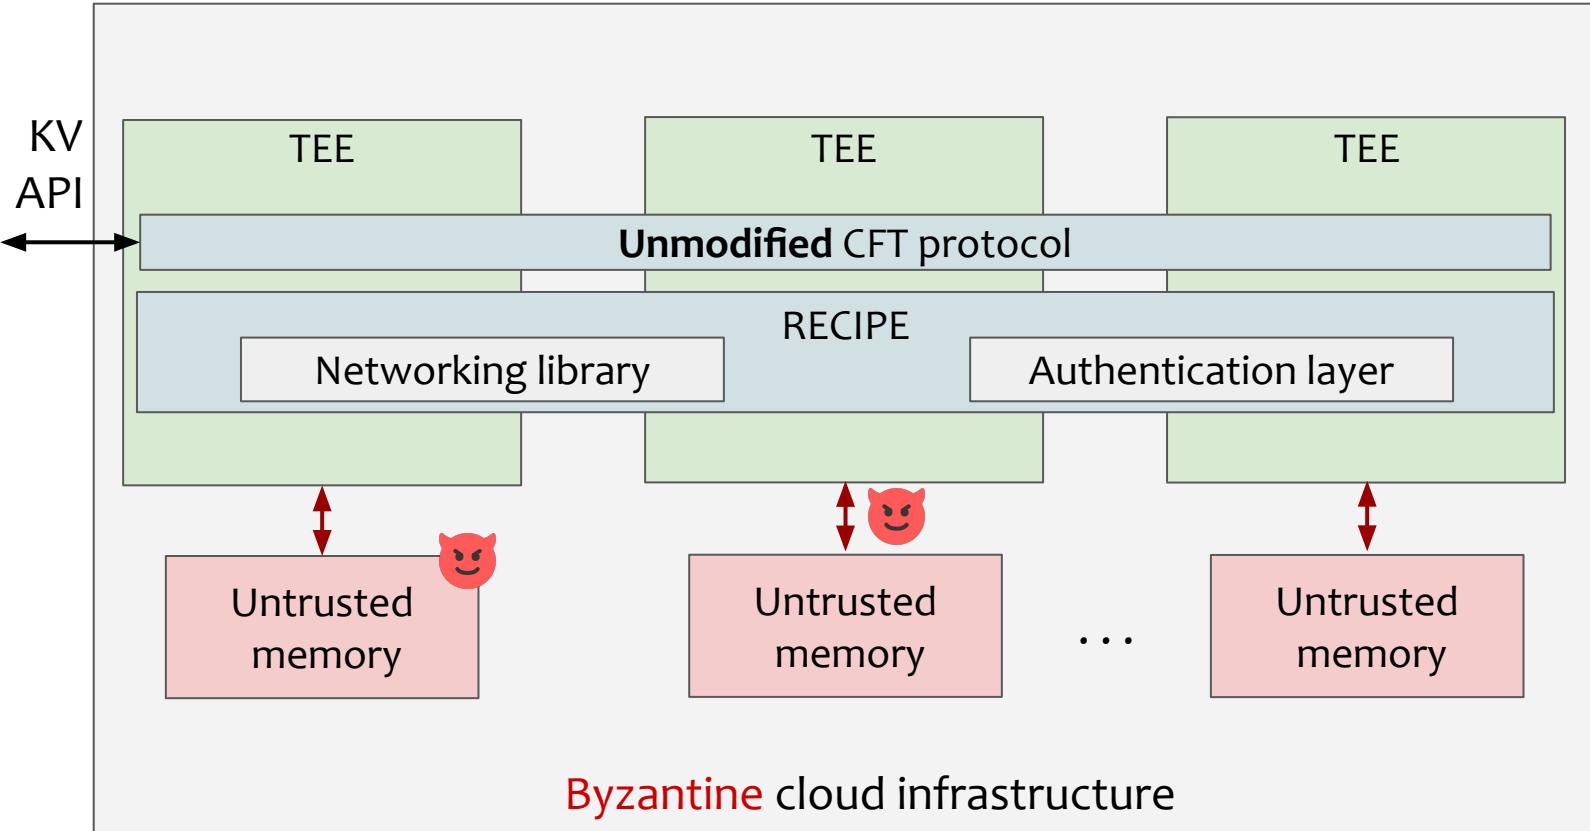
- Robustness
 - can tolerate Byzantine actors in the cloud
- Generality
 - transparent application to a broad category of protocols
- Performance
 - scalability and high-throughput

Outline

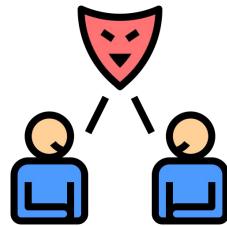
- Motivation

- Overview
- System design
- Evaluation

RECIPE overview



RECIPE ingredients for scalable BFT



#1: Non-equivocation

Do not make conflicting statements
to different nodes



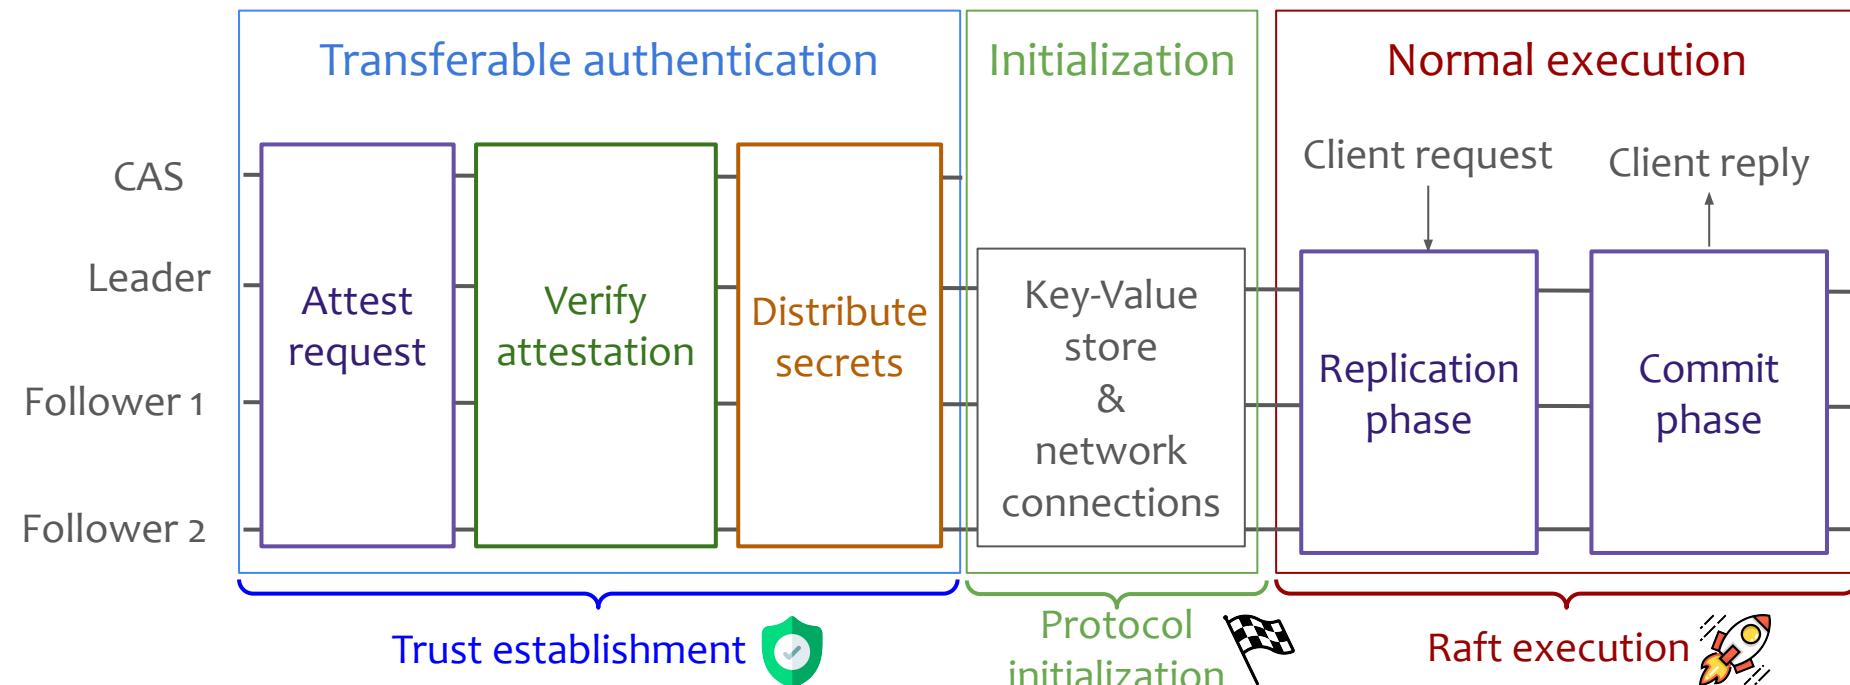
#2: Transferable authentication

Be capable of verifying
the original sender of the message

Allow systems to operate with $2f+1$ nodes in Byzantine environments¹

¹On the (limited) power of non-equivocation, Clement et al., PODC'12.

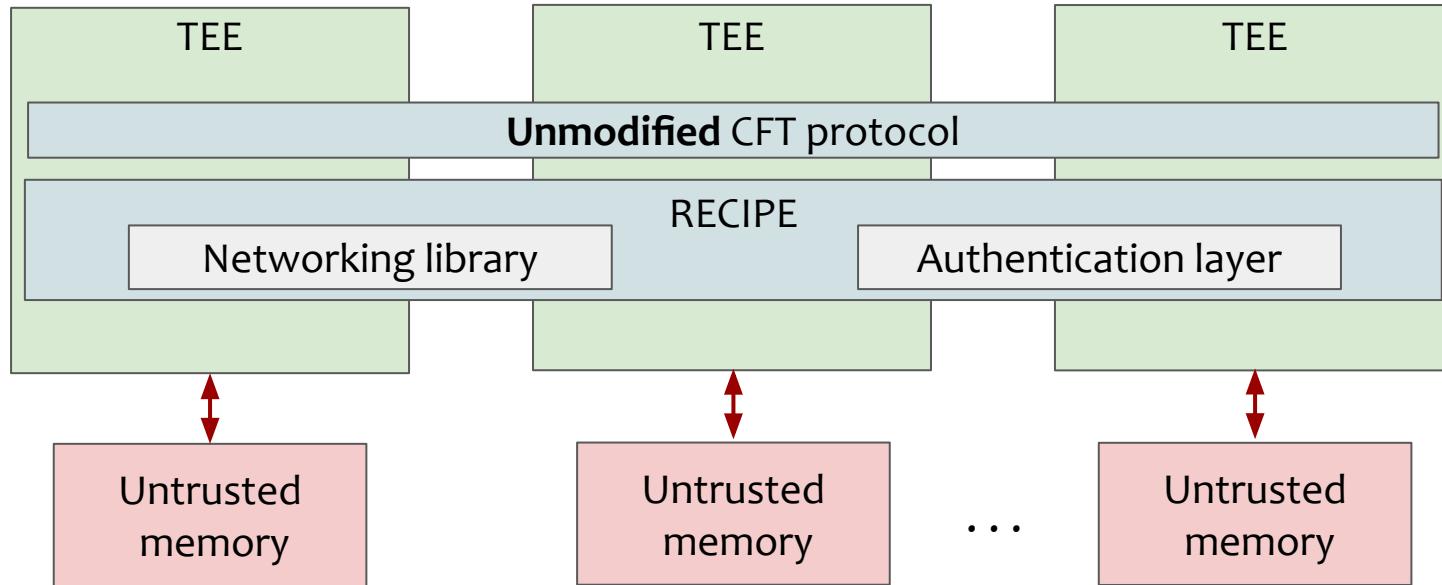
RECIPE protocol: Raft example



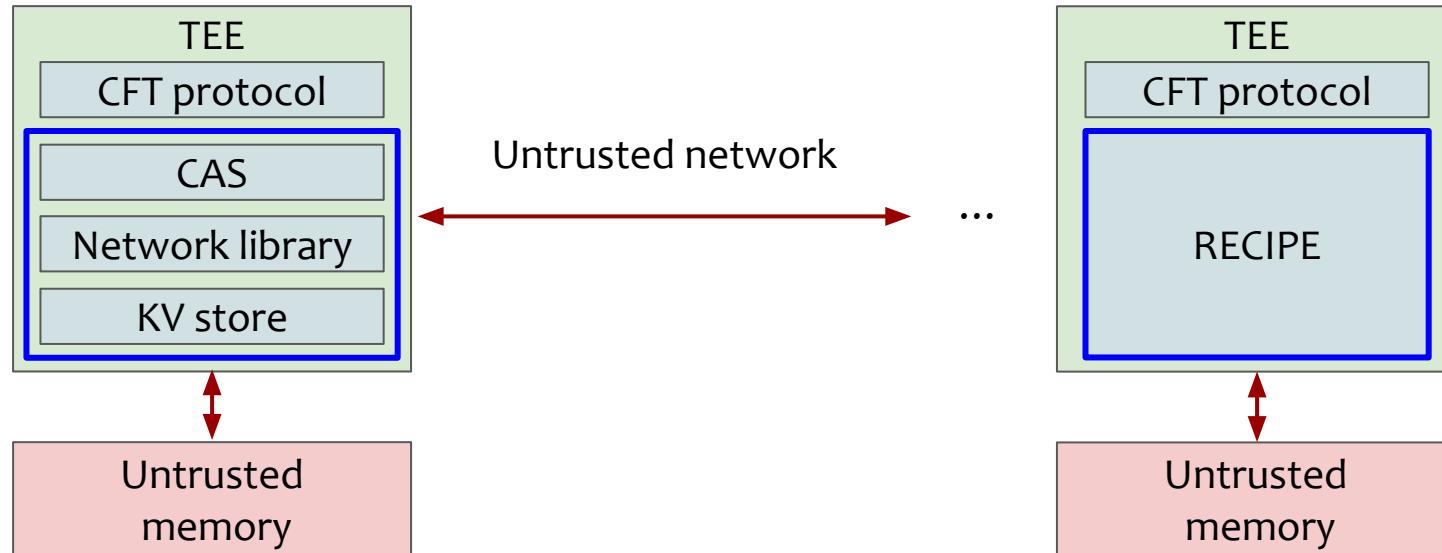
Outline

- Motivation
- Overview
- System design
- Evaluation

A RECIPE node: System stack

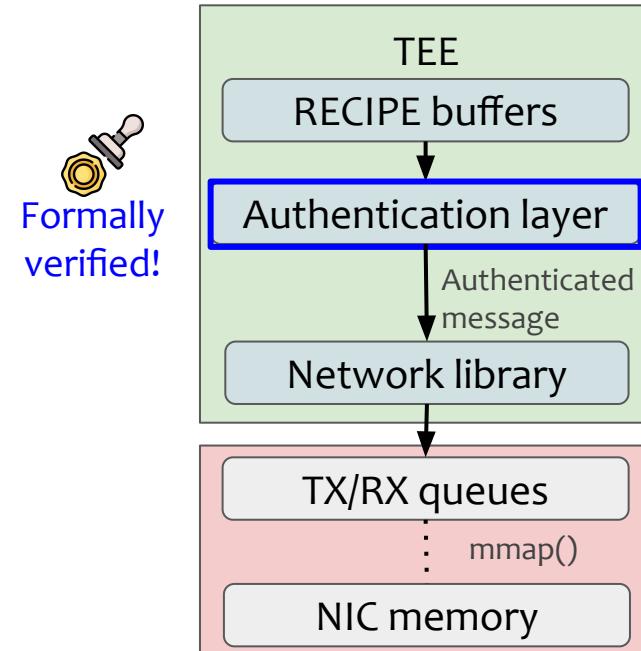


A RECIPE node: System stack



RECIPE network stack and APIs

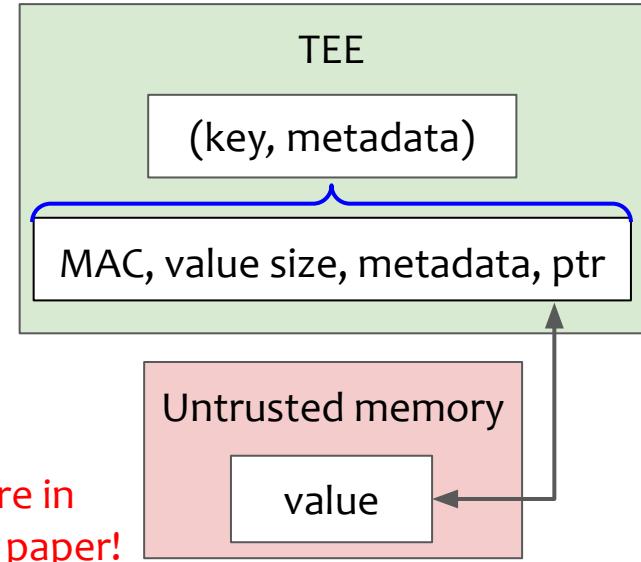
- RECIPE authentication layer
 - transferable authentication, *i.e.*, MACs
 - non-equivocation, *i.e.*, sequence numbers
- RECIPE network library
 - userspace I/O within the TEEs
- RECIPE network APIs
 - authenticated message format
 - RPC-based generic APIs



RECIPE implements user-space **trusted** networking with **generic** APIs

RECIPE Key-Value store

- Overcomes the limited trusted memory
 - splits keys and values
- Maintains the original CFT protocol semantics
 - increases the trust to individual nodes
 - e.g., local (linearizable) reads
- Also, offers confidentiality through encryption
 - not-provided by BFT



More in
the paper!

RECIPE KV store offers **performance** while keeping **the protocol unchanged**

Outline

- Motivation
- Overview
- System design
- Evaluation

Questions:

- What is the performance of RECIPE protocols w.r.t. the state-of-the-art BFT?
- How much overhead do TEEs have in RECIPE?
- What is the performance of RECIPE networking w.r.t. the state-of-the-art?

Experimental setup:

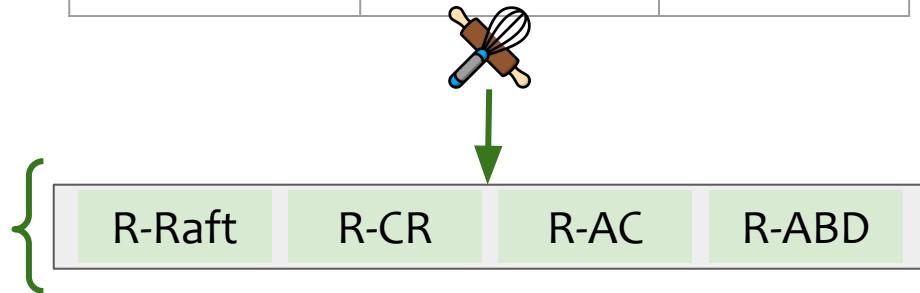
- Distributed system with 3x Intel i9-9900K @3.60GHz. 8 cores
- 40GbE QSFP+ network switch
- YCSB benchmarks, 10k keys, Zipfian distribution

RECIPE application on distributed systems

- We classify CFT protocols
- We select representative protocols
- We transform them with RECIPE

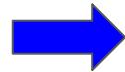
	Leader-based	Leader-less
Total order	Raft [ATC'14]	AC [HPDC'17]
Per-key order	CR [OSDI'04]	ABD [PODC'90]

BFT robustness with CFT performance!

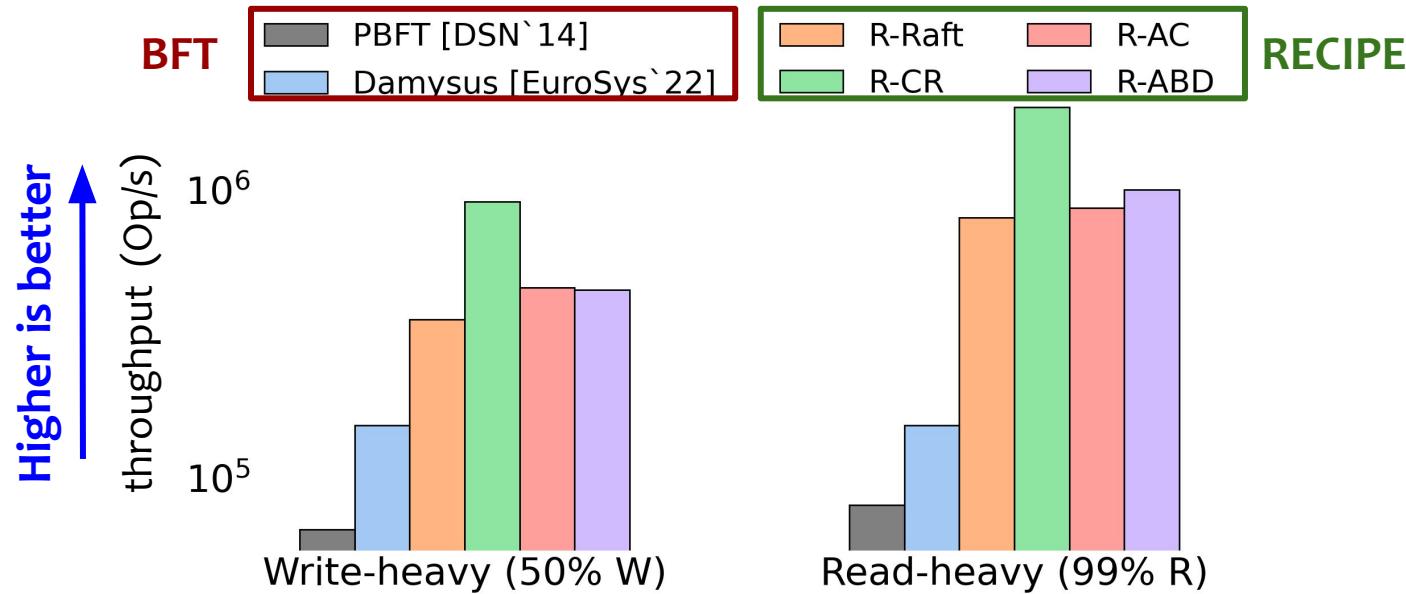


RECIPE is generic and applicable to *any* strongly consistent CFT protocol

Questions:

- 
- What is the performance of RECIPE protocols w.r.t. the state-of-the-art BFT?
 - How much overhead do TEEs have in RECIPE?
 - What is the performance of RECIPE networking w.r.t. the state-of-the-art?

RQ1: RECIPE performance vs. BFT

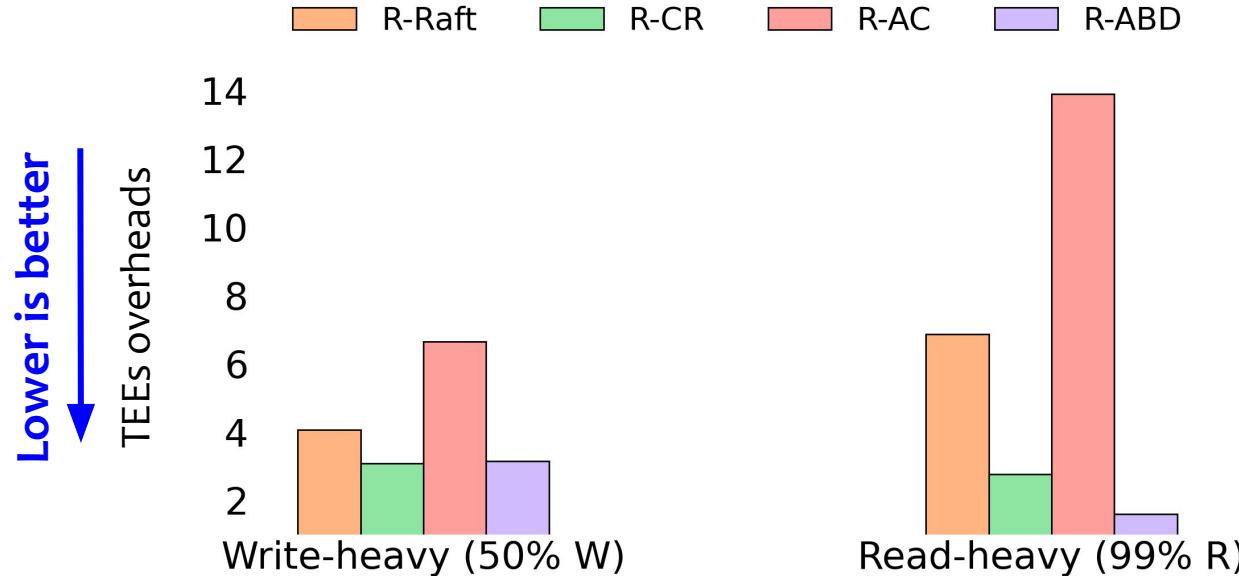


RECIPE achieves **5-20x better performance** compared to state-of-the-art BFT

Questions:

- What is the performance of RECIPE protocols w.r.t. the state-of-the-art BFT?
- How much overhead do TEEs have in RECIPE?
- What is the performance of RECIPE networking w.r.t. the state-of-the-art?

RQ2: TEE overheads

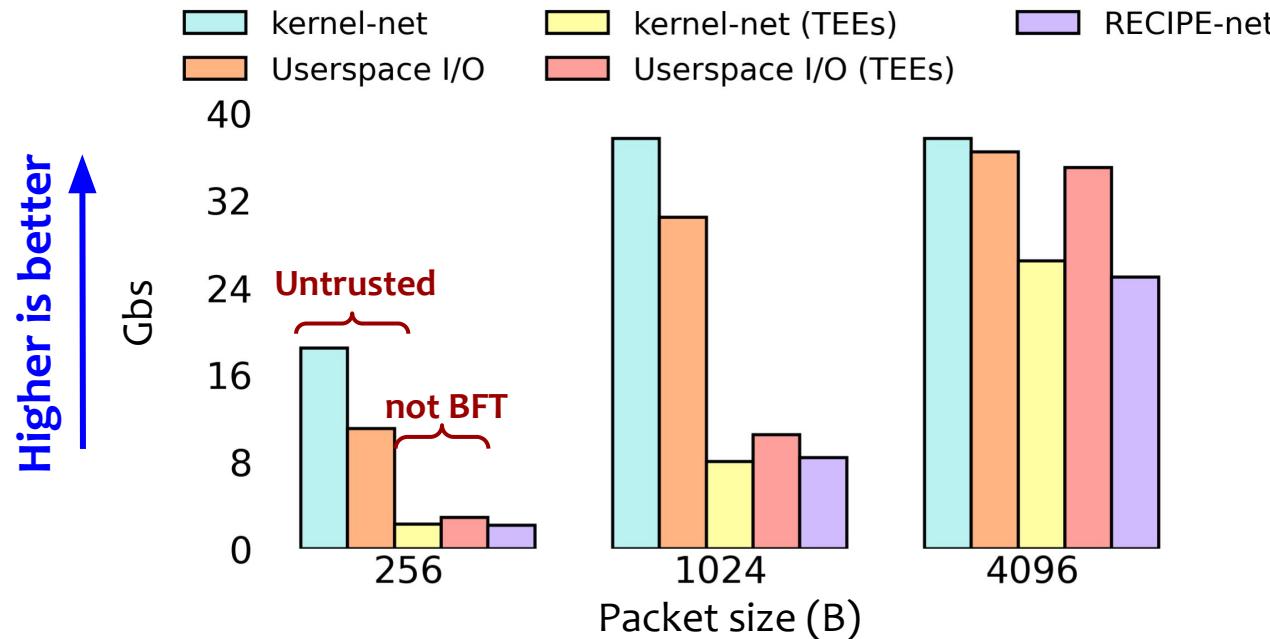


TEEs in RECIPE add **2-14x times slowdown** in throughput

Questions:

- What is the performance of RECIPE protocols w.r.t. the state-of-the-art BFT?
 - How much overhead do TEEs have in RECIPE?
-  What is the performance of RECIPE networking w.r.t. the state-of-the-art?

RQ3: RECIPE's networking performance



RECIPE networking offers **BFT guarantees** while maintaining performance

The CFT vs. BFT conundrum:

- CFT protocols are efficient but **unsuitable** for the **untrusted** cloud
- BFT protocols are robust but **expensive** and **complex**

RECIPE:

- BFT robustness with **high performance** and **scalability**
- **Applicable to any** strongly consistent CFT protocol



Paper



Code